

# nalukanga winnie logistic regression

March 19, 2024

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: ds=pd.read_csv("C:\\Users\\HP\\Desktop\\Social_Network_Ads.csv")
ds
```

```
[2]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

```
[3]: x=ds.drop(['Purchased','Gender'], axis=1)
x
```

```
[3]:
```

	User ID	Age	EstimatedSalary
0	15624510	19	19000
1	15810944	35	20000
2	15668575	26	43000
3	15603246	27	57000
4	15804002	19	76000
..	...	...	...
395	15691863	46	41000
396	15706071	51	23000
397	15654296	50	20000
398	15755018	36	33000
399	15594041	49	36000

[400 rows x 3 columns]

```
[4]: y=ds['Purchased']
      y
```

```
[4]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
     395     1
     396     1
     397     1
     398     0
     399     1
      Name: Purchased, Length: 400, dtype: int64
```

```
[5]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y, test_size=0.
      ↪25,random_state=43)
      x_test.shape
```

```
[5]: (100, 3)
```

```
[6]: from sklearn.linear_model import LogisticRegression
      model=LogisticRegression()
      model
```

```
[6]: LogisticRegression()
```

```
[7]: #checking for missing values
      ds.isnull().sum()
```

```
[7]: User ID      0
      Gender      0
      Age         0
      EstimatedSalary  0
      Purchased    0
      dtype: int64
```

```
[8]: model.fit(x_train,y_train)
```

```
[8]: LogisticRegression()
```

```
[9]: y_pred=model.predict(x_test)
      y_pred
```

```
[9]: array([0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
         0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1,
```

```
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0], dtype=int64)
```

```
[10]: #evaluation
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
```

```
[11]: accuracy_score(y_test, y_pred)
```

```
[11]: 0.71
```

```
[12]: precision_score(y_test, y_pred)
```

```
[12]: 0.7222222222222222
```

```
[13]: recall_score(y_test, y_pred)
```

```
[13]: 0.35135135135135137
```

```
[14]: f1_score(y_test, y_pred)
```

```
[14]: 0.4727272727272727
```

optimization

```
[15]: winnie=LogisticRegression()
winnie
```

```
[15]: LogisticRegression()
```

```
[16]: from sklearn.model_selection import GridSearchCV
param_grid={
    'penalty':['l1', 'l2', 'elasticnet', None],
    'solver':['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag',
↪ 'saga'],
    'C':[1],
    'dual':[True, False]
}
param_grid
```

```
[16]: {'penalty': ['l1', 'l2', 'elasticnet', None],
'solver': ['lbfgs',
'liblinear',
'newton-cg',
'newton-cholesky',
'sag',
'saga'],
'C': [1],
```

```
'dual': [True, False]}
```

```
[17]: grid_search=GridSearchCV(winnie,param_grid,cv=5)
      grid_search
```

```
[17]: GridSearchCV(cv=5, estimator=LogisticRegression(),
                  param_grid={'C': [1], 'dual': [True, False],
                              'penalty': ['l1', 'l2', 'elasticnet', None],
                              'solver': ['lbfgs', 'liblinear', 'newton-cg',
                                         'newton-cholesky', 'sag', 'saga']})
```

```
[18]: import warnings
      warnings.filterwarnings('ignore')
      grid_search.fit(x_train,y_train)
```

```
[18]: GridSearchCV(cv=5, estimator=LogisticRegression(),
                  param_grid={'C': [1], 'dual': [True, False],
                              'penalty': ['l1', 'l2', 'elasticnet', None],
                              'solver': ['lbfgs', 'liblinear', 'newton-cg',
                                         'newton-cholesky', 'sag', 'saga']})
```

```
[19]: best_params= grid_search.best_params_
      best_params
```

```
[19]: {'C': 1, 'dual': False, 'penalty': 'l2', 'solver': 'newton-cholesky'}
```

```
[20]: best_winnie=LogisticRegression(**best_params)
      best_winnie.fit(x_train,y_train)
```

```
[20]: LogisticRegression(C=1, solver='newton-cholesky')
```

```
[21]: y_pred=best_winnie.predict(x_test)
      y_pred
```

```
[21]: array([0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1,
          0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
          0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0], dtype=int64)
```

```
[22]: accuracy_score(y_test,y_pred)
```

```
[22]: 0.78
```

```
[23]: precision_score(y_test,y_pred)
```

```
[23]: 0.7777777777777778
```

```
[24]: recall_score(y_test,y_pred)
```

```
[24]: 0.5675675675675675
```

```
[25]: f1_score(y_test,y_pred)
```

```
[25]: 0.65625
```

```
[ ]:
```